

UCCD2063 ARTIFICIAL INTELLIGENCE TECHNIQUES

Classification of Red Wine Quality using NeuralNetwork

Group ID: 36

Team members:

1. Leong Khei Sheng, 1403785, CS
2. Wan Kar Hou, ID, CS
3. Leong Khai Siang, 1803387, CS

Data Exploration

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

wine = pd.read_csv("./redwinequality.csv")
```

In [2]:

```
wine.head()
```

Out[2]:

	WineID	FixedAcidity	VolatileAcidity	CitricAcid	ResidualSugar	Chlorides	FreeSulfurDioxi
0	1	7.4	0.70	0.00	1.9	0.076	11
1	2	7.8	0.88	0.00	2.6	0.098	25
2	3	7.8	0.76	0.04	2.3	0.092	15
3	4	11.2	0.28	0.56	1.9	0.075	17
4	5	7.4	0.70	0.00	1.9	0.076	11

As per observed from above, 'WineID' is not part of the features and must not included in the model. There are total of 11 features and 'Quality' is the output values.

The features for input matrix have different scales. Standardization on the features is needed to get standard deviation of 1.

In [3]:

```
wine.drop("WineID", axis =1 , inplace = True)
```

In [4]:

```
wine.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
FixedAcidity      1599 non-null float64
VolatileAcidity   1599 non-null float64
CitricAcid        1599 non-null float64
ResidualSugar     1599 non-null float64
Chlorides         1599 non-null float64
FreeSulfurDioxide 1599 non-null float64
TotalSulfurDioxide 1599 non-null float64
Density          1599 non-null float64
PH               1599 non-null float64
Sulphates        1599 non-null float64
Alcohol          1599 non-null float64
Quality          1599 non-null int64
dtypes: float64(11), int64(1)
memory usage: 149.9 KB
```

No empty or null values found in this dataset

In [5]:

```
wine.describe()
```

Out[5]:

	FixedAcidity	VolatileAcidity	CitricAcid	ResidualSugar	Chlorides	FreeSulfurDioxi
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.0000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.8749
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.4601
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.0000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.0000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.0000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.0000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.0000

In [6]:


```
wine.corr()
```

Out[6]:

	FixedAcidity	VolatileAcidity	CitricAcid	ResidualSugar	Chlorides	FreeSul
FixedAcidity	1.000000	-0.256131	0.671703	0.114777	0.093705	
VolatileAcidity	-0.256131	1.000000	-0.552496	0.001918	0.061298	
CitricAcid	0.671703	-0.552496	1.000000	0.143577	0.203823	
ResidualSugar	0.114777	0.001918	0.143577	1.000000	0.055610	
Chlorides	0.093705	0.061298	0.203823	0.055610	1.000000	
FreeSulfurDioxide	-0.153794	-0.010504	-0.060978	0.187049	0.005562	
TotalSulfurDioxide	-0.113181	0.076470	0.035533	0.203028	0.047400	
Density	0.668047	0.022026	0.364947	0.355283	0.200632	
PH	-0.682978	0.234937	-0.541904	-0.085652	-0.265026	
Sulphates	0.183006	-0.260987	0.312770	0.005527	0.371260	
Alcohol	-0.061668	-0.202288	0.109903	0.042075	-0.221141	
Quality	0.124052	-0.390558	0.226373	0.013732	-0.128907	

In [7]:

```
wine.hist(bins=50,figsize=(20,15))
plt.show()
```



In [8]:

```
wine.isnull().sum()
```

Out[8]:

```
FixedAcidity      0
VolatileAcidity   0
CitricAcid        0
ResidualSugar     0
Chlorides         0
FreeSulfurDioxide 0
TotalSulfurDioxide 0
Density          0
PH               0
Sulphates         0
Alcohol          0
Quality          0
dtype: int64
```


In [9]:

```
print(wine.Quality.value_counts())
```

```
5    681
6    638
7    199
4     53
8     18
3     10
Name: Quality, dtype: int64
```

In [10]:

```
sns.countplot(x="Quality", data=wine)
plt.show()
```



The output values is skewed(right).

Preprocessing

To obtain a more accurate data, the output value is divided into 2 categories which 3-5 are categorized as 0 (Low Quality), and 1 (High Quality)

In [11]:

```
wine['Quality'] = pd.cut(wine['Quality'], bins=2, labels=False)
```

In [12]:

```
X = wine.drop('Quality', axis = 1)  
y = wine["Quality"]
```

In [13]:

```
y.value_counts()
```

Out[13]:


```
1    855  
0    744  
Name: Quality, dtype: int64
```

In [14]:

```
sns.countplot(y)
```

Out[14]:

<matplotlib.axes._subplots.AxesSubplot at 0xbc90a30>



The output values are more balanced after converted to 0 and 1

In [15]:

```
print ('X: shape=', X.shape)  
print ('y: shape=', y.shape)
```

```
X: shape= (1599, 11)  
y: shape= (1599,)
```

The datasets are split into 80% training set and 20% test set.

In [16]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
print ('train shape=', X_train.shape)
print ('test shape=', X_test.shape)
```

train shape= (1279, 11)

test shape= (320, 11)

Standardization on the features needed to be performed as the features have different scales.

In [17]:

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler(copy = False)
scaler.fit(X_train)
X_train_tr = scaler.transform(X_train)
y_train = y_train.values
```

Training

Neural Network

Create the model using MLPClassifier in Neural Network.

In [18]:

```
from sklearn.neural_network import MLPClassifier

mlp = MLPClassifier(random_state=42)
mlp.fit(X_train_tr, y_train)
y_pred = mlp.predict(X_train_tr)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
  % self.max_iter, ConvergenceWarning)
```

In [19]:

```
def peek_results(actual, predicted, num = 20):
    print('Actual | Predicted')
    print('-----')
    for i in range(num):
        sel = np.random.randint(0, len(y_train))
        print(actual[sel], ' | ', predicted[sel])

peek_results(y_train, y_pred)
```

Actual | Predicted

```
-----
0      |      0
1      |      1
0      |      0
1      |      1
1      |      1
0      |      0
1      |      0
0      |      1
1      |      1
1      |      1
1      |      1
1      |      1
0      |      0
1      |      1
0      |      0
1      |      1
0      |      0
0      |      0
1      |      1
1      |      1
```

Validation

Import all the libraries that required later.

In [20]:

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.model_selection import cross_validate
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score, recall_score, f1_score
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
```

PCA

Principal Component Analysis (PCA) is used to reduce the number of features of the data. Some output of the dataset is affected by only few main features, thus PCA is applied in this datasets to test whether reducing number of fetures will increase the accuracy of data.

In [21]:

```
from sklearn.decomposition import PCA

pca = PCA(n_components=3)
pca.fit(X_train_tr, y_train)
X_train_tr_pca = pca.transform(X_train_tr)

mlp.fit(X_train_tr_pca, y_train)
pca_pred = mlp.predict(X_train_tr_pca)
print('Training Accuracy with PCA = ', accuracy_score(y_train, pca_pred))
```

Training Accuracy with PCA = 0.7498045347928068

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
% self.max_iter, ConvergenceWarning)

The training accuracy is around 75% after reducing to 3 features for the data.

Accuracy

In [22]:

```
train_acc = accuracy_score(y_train, y_pred)

print("Training accuracy without PCA: {:.4f}".format(train_acc))
```

Training accuracy without PCA: 0.8139

The training accuracy without applying PCA is around 81%. Thus all features in this datasets are important and contribute much to the output values.

The model is fitted again with the original dataset.

In [23]:

```
mlp.fit(X_train_tr, y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
% self.max_iter, ConvergenceWarning)

Out[23]:

```
MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=
0.9,
              beta_2=0.999, early_stopping=False, epsilon=1e-08,
              hidden_layer_sizes=(100,), learning_rate='constant',
              learning_rate_init=0.001, max_iter=200, momentum=0.9,
              n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
              random_state=42, shuffle=True, solver='adam', tol=0.0001,
              validation_fraction=0.1, verbose=False, warm_start=False)
```


Cross Validation

The training accuracy is measured with cross-validation to evaluate the model. 4-folds is set for this model.

In [24]:

```
kfold_result = cross_validate (mlp, X_train_tr, y_train, cv=4, scoring='accuracy', return_train_score=True)
```

```
print('Average training accuracy = {:.4f}'.format(kfold_result['train_score'].mean()))
print('Average validation accuracy = {:.4f}'.format(kfold_result['test_score'].mean()))
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
```

```
% self.max_iter, ConvergenceWarning)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
```

```
% self.max_iter, ConvergenceWarning)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
```

```
% self.max_iter, ConvergenceWarning)
```

```
Average training accuracy = 0.8108
```

```
Average validation accuracy = 0.7577
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
```

```
% self.max_iter, ConvergenceWarning)
```

Confusion Matrix

In [25]:

```
y_pred_cv = cross_val_predict (mlp, X_train_tr, y_train, cv=4)
```

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix (y_train, y_pred_cv)
print(cm)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

```
[[456 147]
 [163 513]]
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

Precision, Recall, F1 Score

In [26]:

```
print(classification_report(y_train, y_pred_cv))
```

	precision	recall	f1-score	support
0	0.74	0.76	0.75	603
1	0.78	0.76	0.77	676
micro avg	0.76	0.76	0.76	1279
macro avg	0.76	0.76	0.76	1279
weighted avg	0.76	0.76	0.76	1279

In [27]:

```
print('Precision = {:.4f}'.format(precision_score(y_train, y_pred_cv)))
print('Recall = {:.4f}'.format(recall_score (y_train, y_pred_cv)))
print('F1 Score = {:.4f}'.format(f1_score(y_train, y_pred_cv)))
```

Precision = 0.7773

Recall = 0.7589

F1 Score = 0.7680

In [28]:

```
kfold_scores = cross_validate (mlp, X_train_tr, y_train, cv=4, scoring = ['precision',
'recall', 'f1'], return_train_score=False)

print('Average cross-validation recall = {:.4f}'.format(kfold_scores['test_recall'].mean()))
print('Average cross-validation precision = {:.4f}'.format(kfold_scores['test_precision'].mean()))
print('Average cross-validation f1 = {:.4f}'.format(kfold_scores['test_f1'].mean()))
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

Average cross-validation recall = 0.7589

Average cross-validation precision = 0.7775

Average cross-validation f1 = 0.7679

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

Precision-Recal Graph

In [29]:

```
y_scores_cv = cross_val_predict (mlp, X_train_tr, y_train, cv=4)
precisions, recalls, thresholds = precision_recall_curve (y_train, y_scores_cv)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

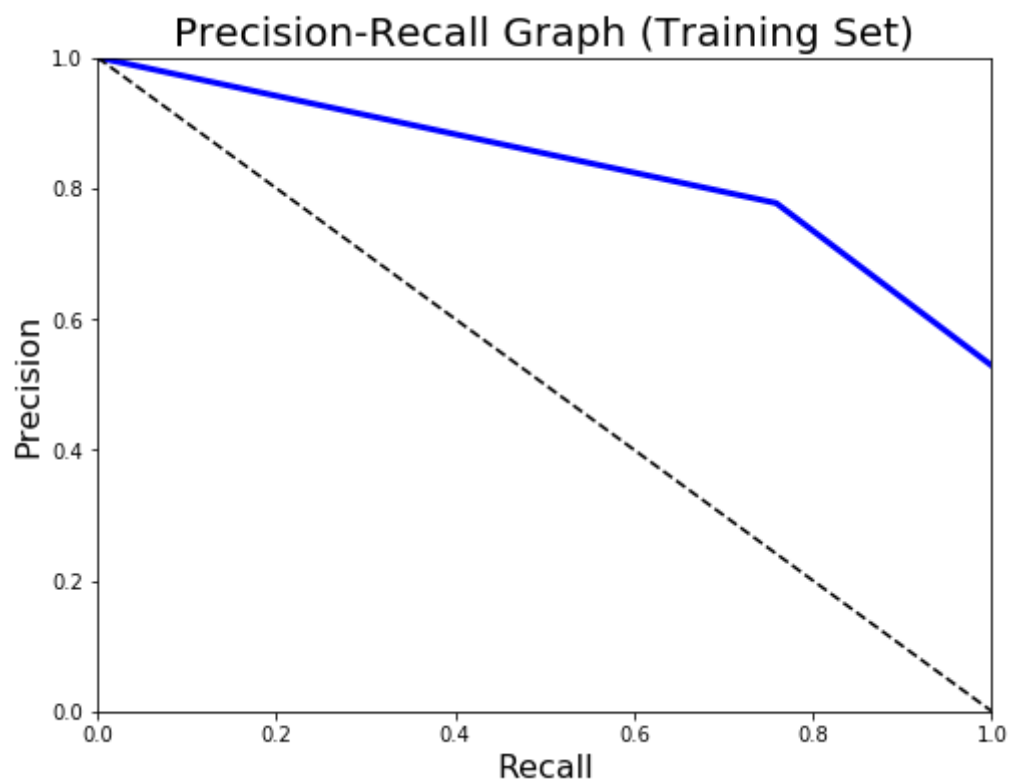
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:562: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

% self.max_iter, ConvergenceWarning)

In [30]:

```
def plot_precision_vs_recall(precisions, recalls):
    plt.plot(recalls, precisions, "b-", linewidth=3)
    plt.plot(np.linspace(0, 1, 20), np.linspace(1, 0, 20), 'k--')
    plt.xlabel("Recall", fontsize=16)
    plt.ylabel("Precision", fontsize=16)
    plt.axis([0, 1, 0, 1])

plt.figure(figsize=(8, 6))
plot_precision_vs_recall(precisions, recalls)
plt.title('Precision-Recall Graph (Training Set)', fontsize = 20)
plt.show()
```

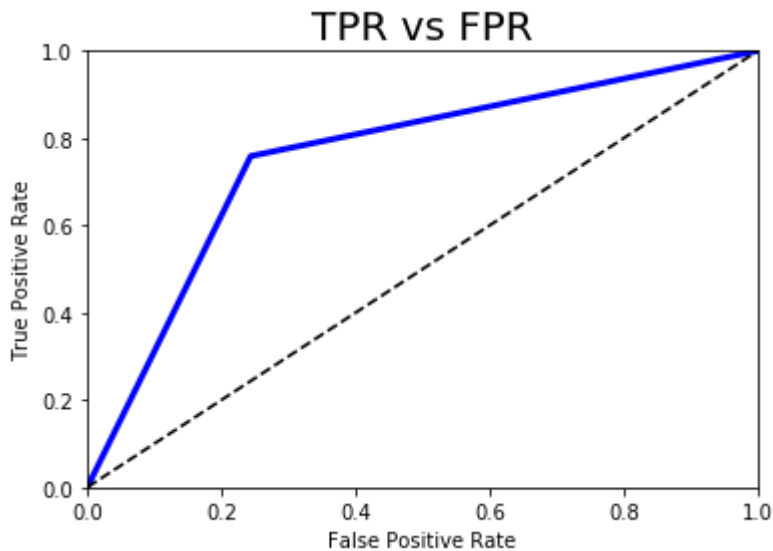


In [31]:

```
fpr, tpr, thresholds = roc_curve(y_train, y_scores_cv)

def plot_roc_curve (fpr, tpr, style = 'b-', label = None):
    plt.plot(fpr, tpr, style, linewidth = 3, label = label)
    plt.plot([0,1], [0, 1], 'k--')
    plt.axis([0, 1, 0, 1])
    plt.xlabel ('False Positive Rate')
    plt.ylabel ('True Positive Rate')
    plt.title('TPR vs FPR', fontsize=20)

plot_roc_curve(fpr, tpr)
```



In [32]:

```
auc = roc_auc_score(y_train, y_scores_cv)
print('AUC = {:.4f}'.format(auc))
```

AUC = 0.7575

Fine Tune

GridSearchCV

To get the best hyperparameters, GridSearchCV is used.

In [33]:

```
import time
from sklearn.model_selection import GridSearchCV

start_time = time.time()
param = {'solver':['lbfgs'], 'max_iter':[500,1000,1500], 'alpha':10.0**-np.arange(1,10), 'hidden_layer_sizes':np.arange(10,15)}
grid_search = GridSearchCV(mlp, param, cv=4, scoring='accuracy')
print('Performing grid search...', end = '')
grid_search.fit(X_train_tr, y_train)
print('Completed', end = ' ')
print('{:.2f}s'.format(time.time()-start_time))
```

Performing grid search...Completed 444.38s

In [34]:

```
best_model = grid_search.best_estimator_
best_model
```

Out[34]:

```
MLPClassifier(activation='relu', alpha=0.1, batch_size='auto', beta_1=0.9,
              beta_2=0.999, early_stopping=False, epsilon=1e-08,
              hidden_layer_sizes=12, learning_rate='constant',
              learning_rate_init=0.001, max_iter=1000, momentum=0.9,
              n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
              random_state=42, shuffle=True, solver='lbfgs', tol=0.0001,
              validation_fraction=0.1, verbose=False, warm_start=False)
```

In [35]:

```
print('Accuracy of the best model =', grid_search.best_score_)
```

Accuracy of the best model = 0.7505863956215794

Testing

In [36]:

```
X_test_tr = scaler.transform(X_test)
y_test = y_test.values
```

Use the best model to predict test data.

In [37]:

```
y_pred_test = best_model.predict(X_test_tr)
print('Test data Accuracy= ', accuracy_score(y_test, y_pred_test))
```

Test data Accuracy= 0.778125

In [38]:

```
print('Test precision = {:.4f}'.format(precision_score(y_test, y_pred_test)))
print('Test recall = {:.4f}'.format(recall_score (y_test, y_pred_test)))
print('Test f1 score = {:.4f}'.format(f1_score(y_test, y_pred_test)))
```

Test precision = 0.8068

Test recall = 0.7933

Test f1 score = 0.8000

Compare with Other Classifiers

RandomForestClassifier

In [39]:

```
from sklearn.ensemble import RandomForestClassifier

forest = RandomForestClassifier(n_estimators=10, random_state=42)
forest.fit(X_train_tr, y_train)
y_pred_forest = forest.predict(X_train_tr)

print('Accuracy = ', accuracy_score(y_train, y_pred_forest))
```

Accuracy = 0.9929632525410477

The model fitted very well on the training set with accuracy 99%.

In [40]:

```
print(classification_report(y_train, y_pred_forest))
```

	precision	recall	f1-score	support
0	0.99	1.00	0.99	603
1	1.00	0.99	0.99	676
micro avg	0.99	0.99	0.99	1279
macro avg	0.99	0.99	0.99	1279
weighted avg	0.99	0.99	0.99	1279

In [41]:

```
y_pred_test_forest = forest.predict(X_test_tr)

print('Test data Accuracy= ', accuracy_score(y_test, y_pred_test_forest))
```

Test data Accuracy= 0.778125

However, the model doesn't fitted so well on the test set. This may caused by the model is slightly overfitting.

In [42]:

```
kfold_result = cross_validate (forest, X_train_tr, y_train, cv=5, scoring='accuracy', return_train_score=True)
print('Average training accuracy = {:.4f}'.format(kfold_result['train_score'].mean()))
print('Average validation accuracy = {:.4f}'.format(kfold_result['test_score'].mean()))
```

Average training accuracy = 0.9893
Average validation accuracy = 0.7764

RandomForestClassifier seem like perform better than MLPClassifier. However, the model is slightly overfitting the datasets.

SVCClassifier

In [43]:

```
from sklearn.svm import SVC

svc = SVC()
svc.fit(X_train_tr, y_train)
y_pred_svc = svc.predict(X_train_tr)

print('Accuracy = ', accuracy_score(y_train, y_pred))
```

Accuracy = 0.8139171227521501

In [44]:

```
print(classification_report(y_train, y_pred_svc))
```

	precision	recall	f1-score	support
0	0.78	0.81	0.79	603
1	0.82	0.80	0.81	676
micro avg	0.80	0.80	0.80	1279
macro avg	0.80	0.80	0.80	1279
weighted avg	0.80	0.80	0.80	1279

In [45]:

```
y_pred_test_svc = svc.predict(X_test_tr)

print('Accuracy = ', accuracy_score(y_test, y_pred_test_svc))
```

Accuracy = 0.771875

In [46]:

```
kfold_result = cross_validate (svc, X_train_tr, y_train, cv=5, scoring='accuracy', return_train_score=True)
print('Average training accuracy = {:.4f}'.format(kfold_result['train_score'].mean()))
print('Average validation accuracy = {:.4f}'.format(kfold_result['test_score'].mean()))
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

Average training accuracy = 0.8030

Average validation accuracy = 0.7546

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

The training accuracy of SVCClassifier is quite same as MLPClassifier. MLPClassifier have slightly better performace during prediction of test data.